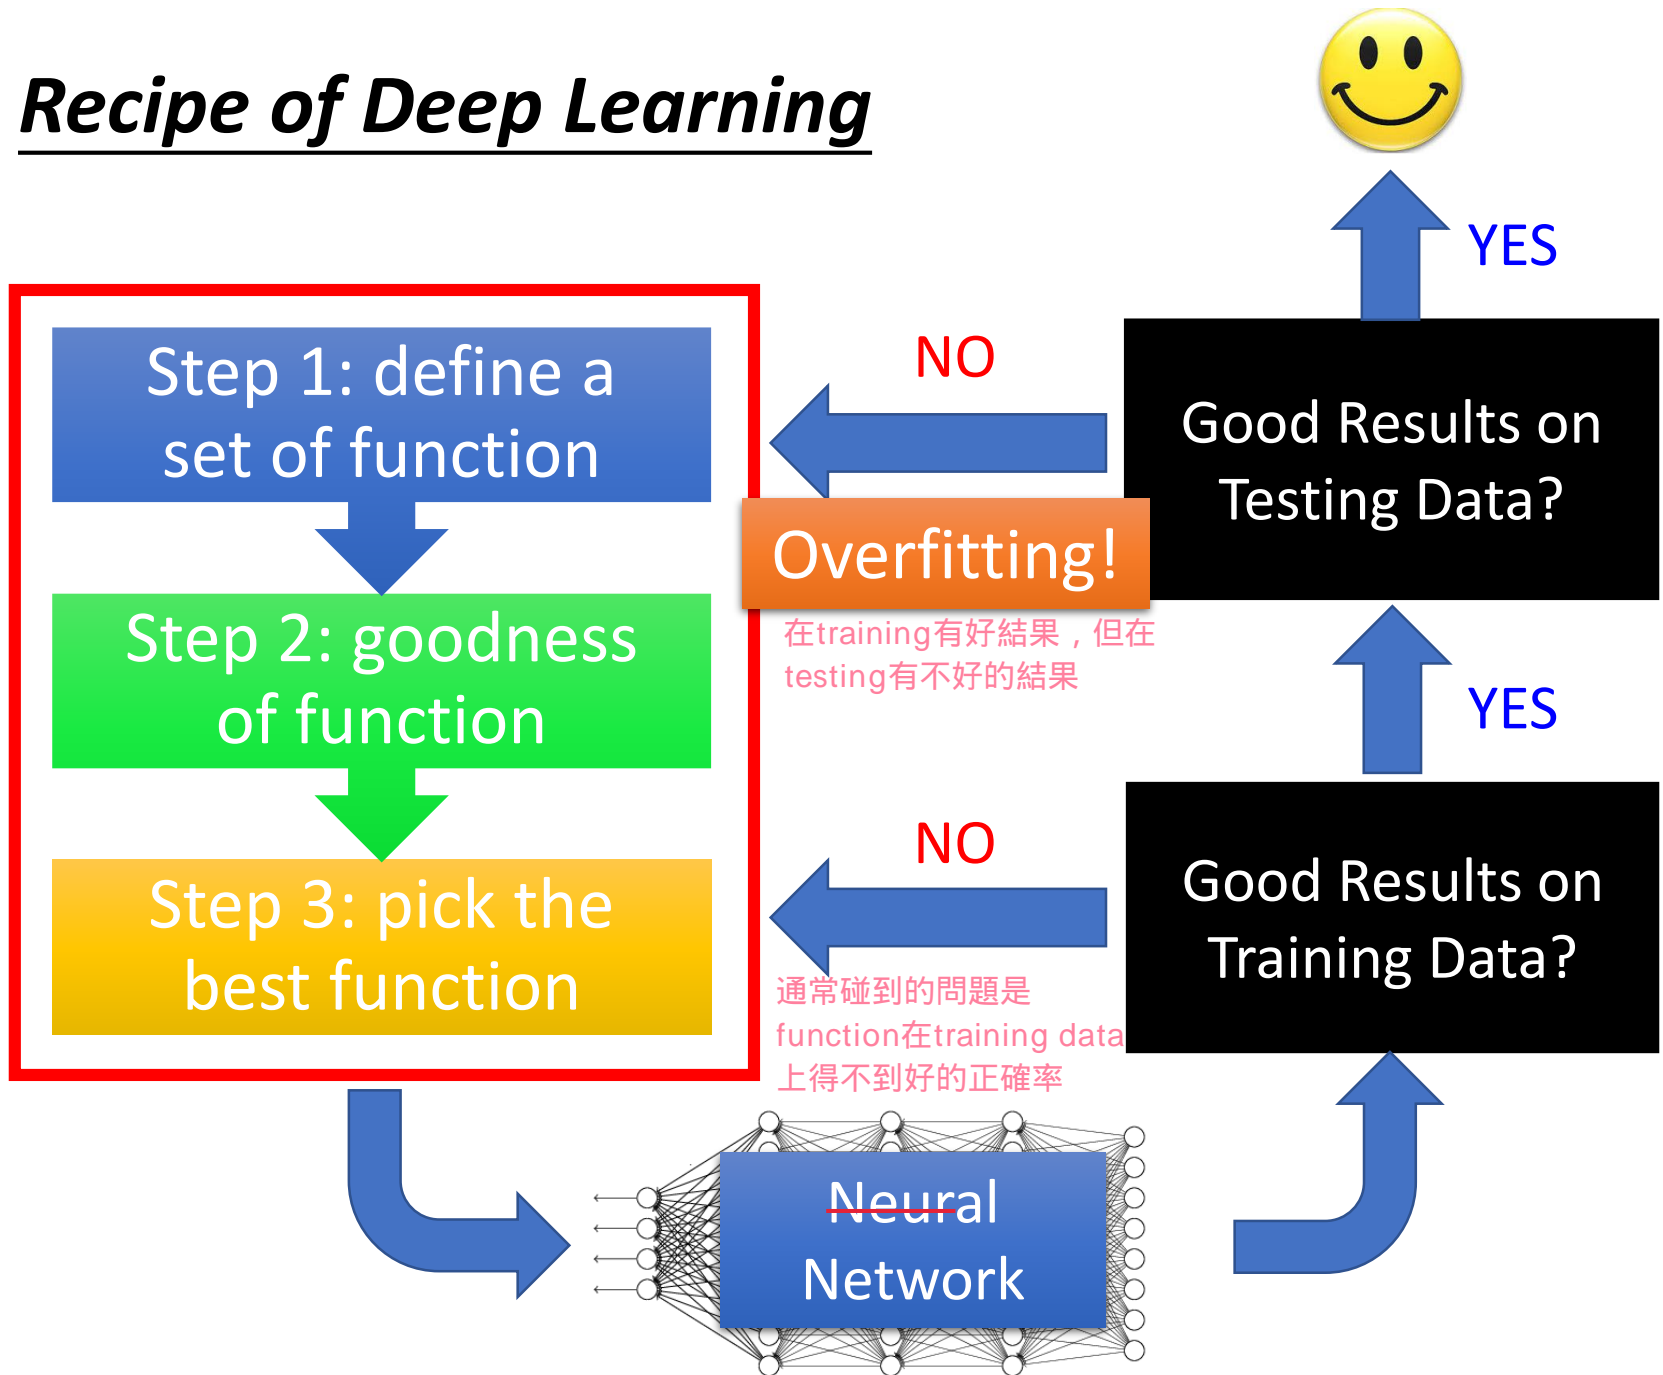



# Tips for Deep Learning

# Recipe of Deep Learning



# Do not always blame Overfitting



Deep Residual Learning for Image Recognition  
<http://arxiv.org/abs/1512.03385>



# Recipe of Deep Learning

思考這個方法是要解什麼問題

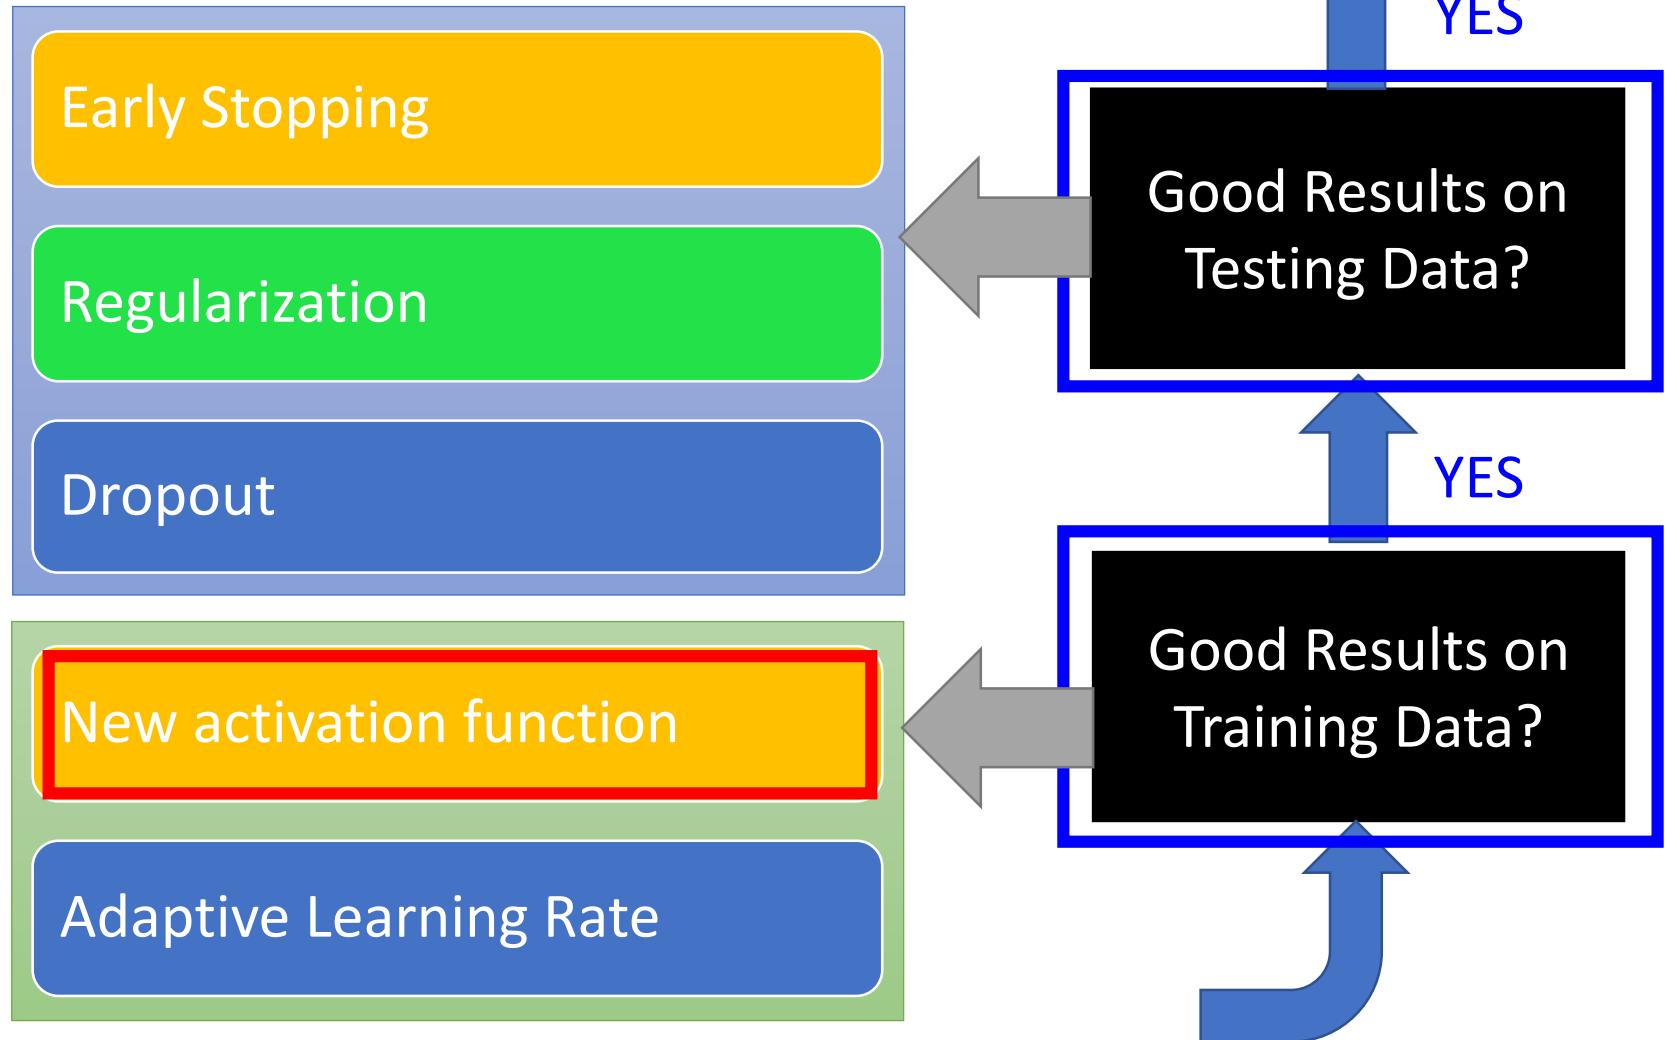
Different approaches for different problems.

e.g. dropout for good results on testing data

testing結果不好時才應使用dropout




# Recipe of Deep Learning




架構不好?

# Hard to get the power of Deep ...



# Vanishing Gradient Problem



靠近input的微分值較小

Smaller gradients

Learn very slow

Almost random

若設定相同的  
learning  
rate

Larger gradients


Learn very fast

Already converge


靠近output的微分值較大

based on random!?

# Vanishing Gradient Problem



影響會衰減  
所以通過越多層，最後對  
output 的影響越小



eg. sigmoid function 會把值硬壓在 0 跟 1 之間

Intuitive way to compute the derivatives ...


$$\frac{\partial l}{\partial w} = ? \quad \frac{\Delta l}{\Delta w}$$



# ReLU

改activation function , 現在常用ReLU

- Rectified Linear Unit (ReLU)




[Xavier Glorot, AISTATS'11]  
[Andrew L. Maas, ICML'13]  
[Kaiming He, arXiv'15]

## Reason:

1. Fast to compute
2. Biological reason
3. Infinite sigmoid with different biases
4. Vanishing gradient problem

ReLU等於無窮  
多的sigmoid疊  
加的結果


# ReLU



0的neuron沒影響，可拿掉

# ReLU


A Thinner linear network




因為ReLU的input=output，所以不會有前面提到影響遞減的問題

# ReLU - variant

*Leaky ReLU*



*Parametric ReLU*



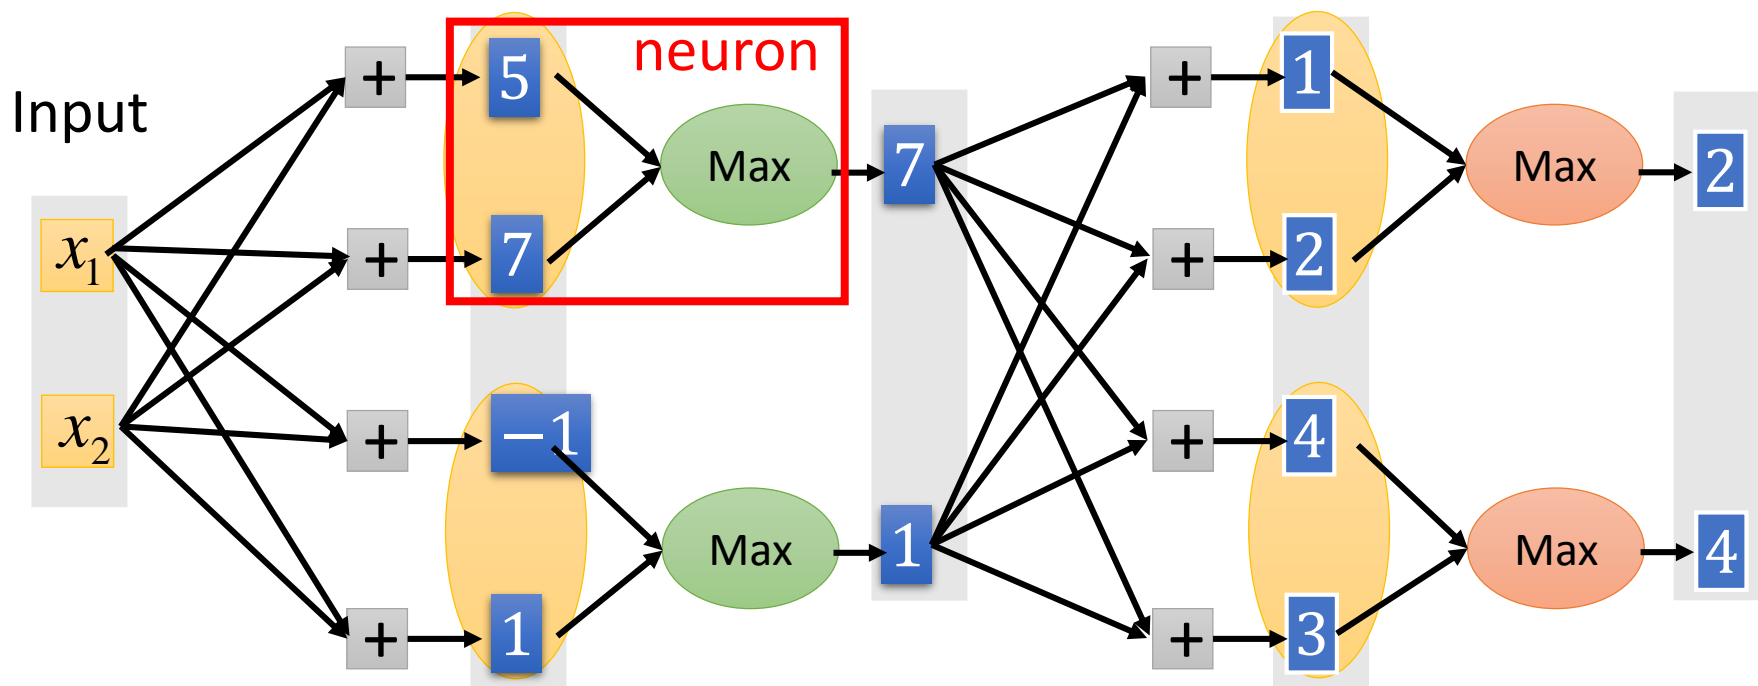
$\alpha$  also learned by  
gradient descent

# Maxout

ReLU is a special cases of Maxout

機器自動學

- Learnable activation function [Ian J. Goodfellow, ICML'13]







You can have more than 2 elements in a group.

手動設幾個element放一組




# Maxout

ReLU is a special cases of Maxout




# Maxout

More than ReLU




Learnable Activation Function




# Maxout

- Learnable activation function [\[Ian J. Goodfellow, ICML'13\]](#)
  - Activation function in maxout network can be any piecewise linear convex function
  - How many pieces depending on how many elements in a group

2 elements in a group



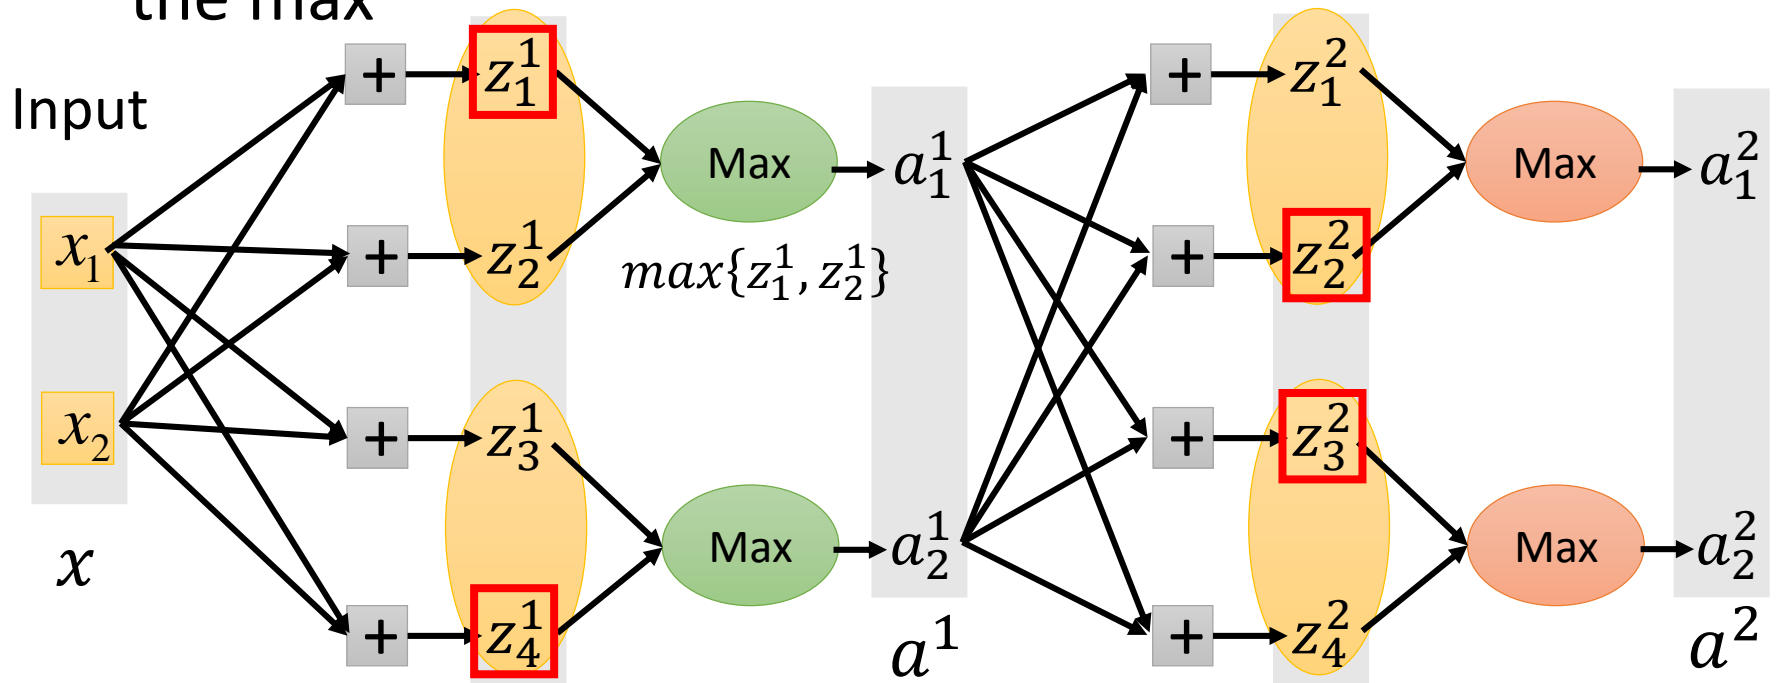
3 elements in a group






# Maxout - Training

- Given a training data  $x$ , we know which  $z$  would be the max



# Maxout - Training

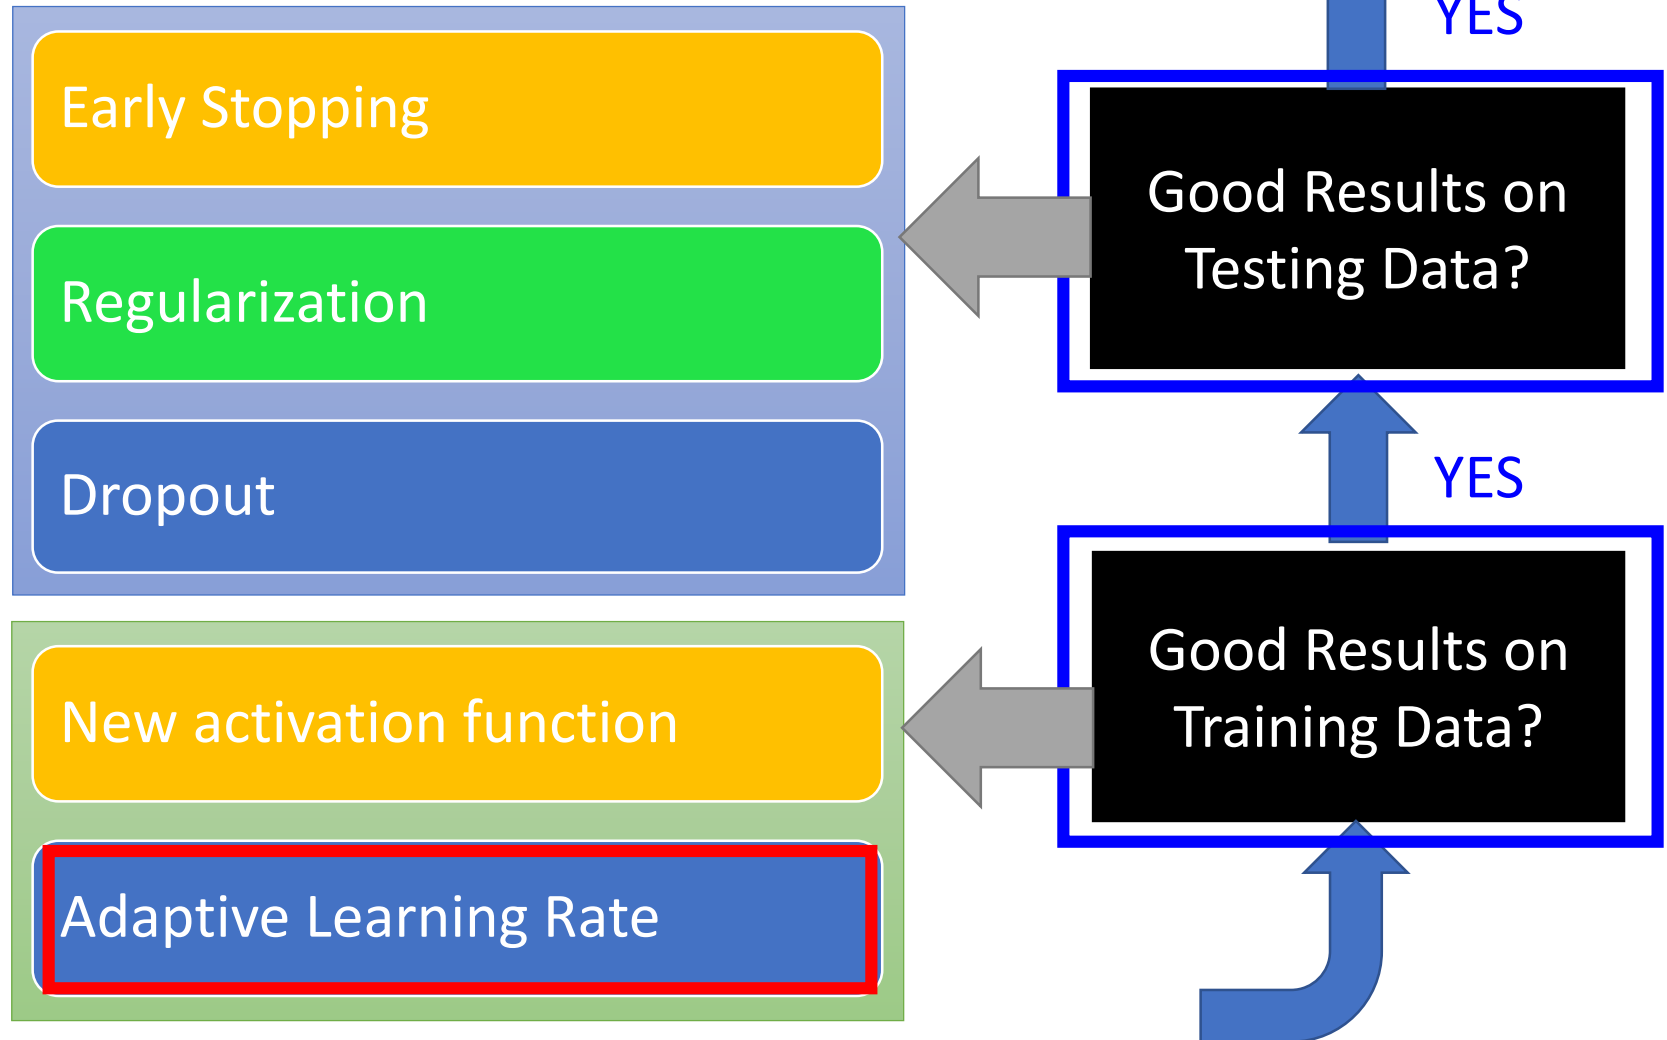
- Given a training data  $x$ , we know which  $z$  would be the max




- Train this thin and linear network

Different thin and linear network for different examples

# Recipe of Deep Learning



# Review




## Adagrad

$$w^{t+1} \leftarrow w^t - \frac{\eta}{\sqrt{\sum_{i=0}^t (g^i)^2}} g^t$$

Use first derivative to estimate second derivative

# RMSProp

Error Surface can be very complex when training NN.



# RMSProp

$$w^1 \leftarrow w^0 - \frac{\eta}{\sigma^0} g^0 \quad \sigma^0 = g^0$$

$$w^2 \leftarrow w^1 - \frac{\eta}{\sigma^1} g^1 \quad \sigma^1 = \sqrt{\alpha(\sigma^0)^2 + (1 - \alpha)(g^1)^2}$$

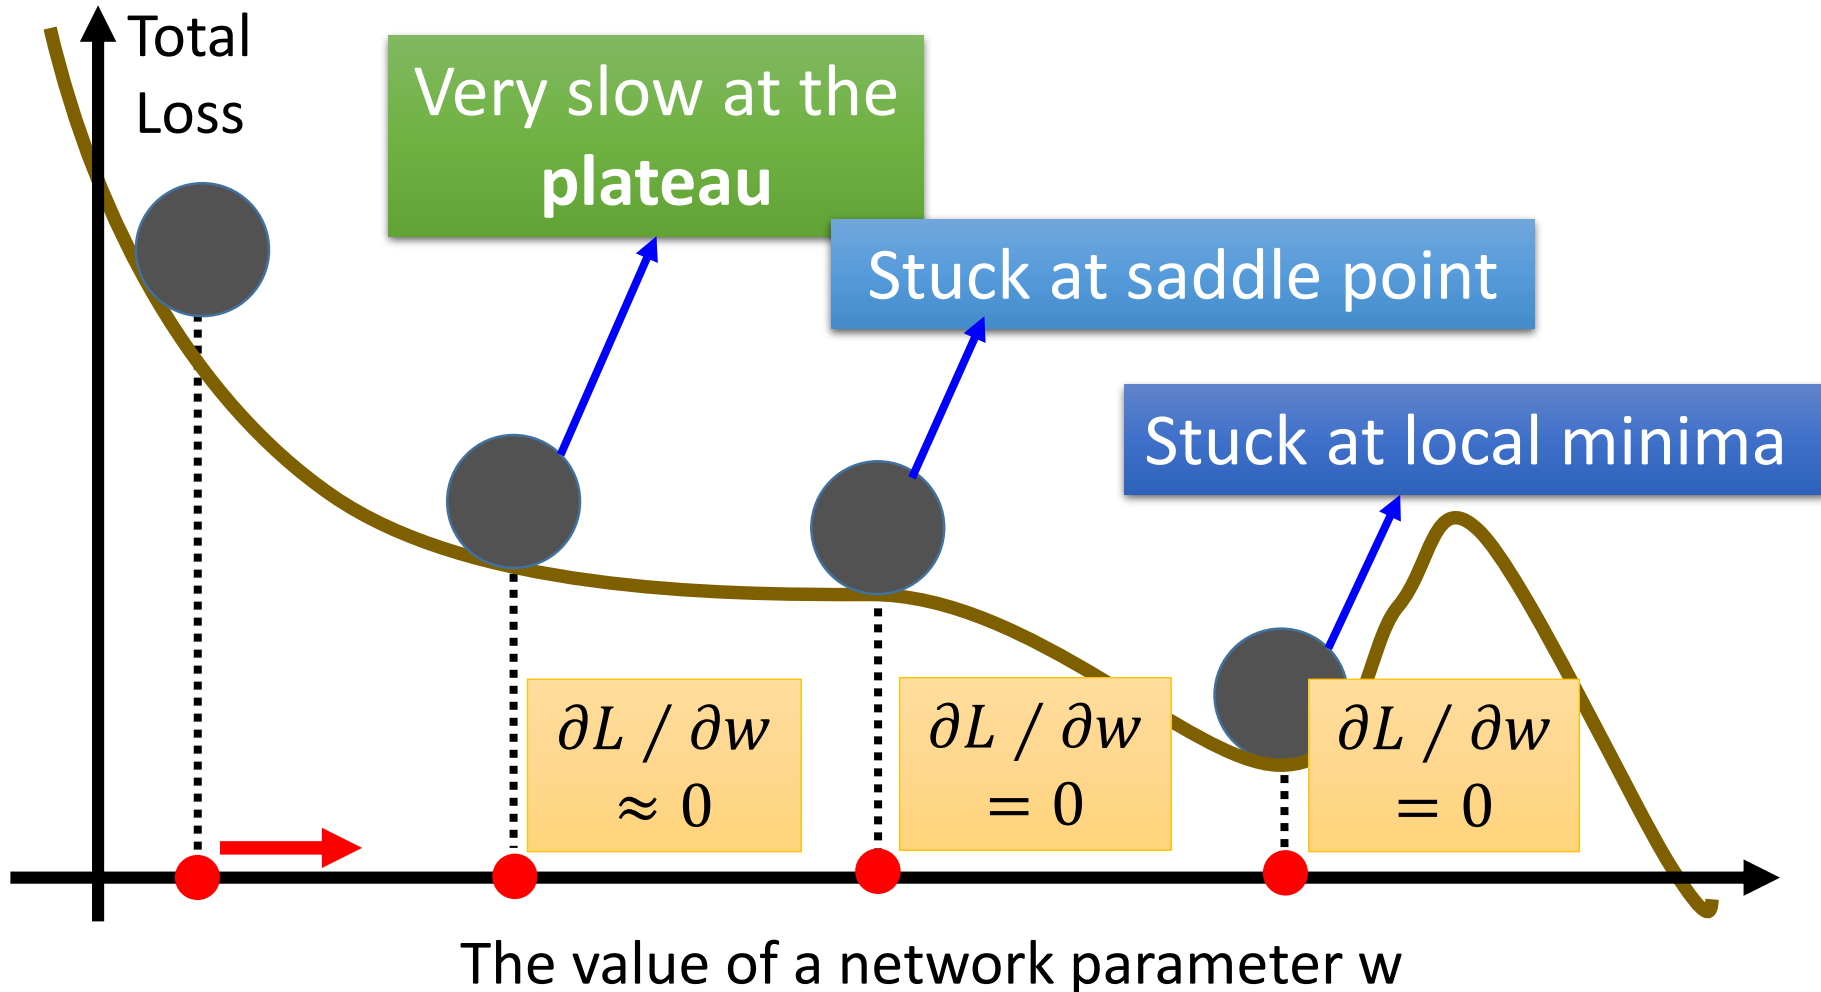
$$w^3 \leftarrow w^2 - \frac{\eta}{\sigma^2} g^2 \quad \sigma^2 = \sqrt{\alpha(\sigma^1)^2 + (1 - \alpha)(g^2)^2}$$

$\vdots$

$$w^{t+1} \leftarrow w^t - \frac{\eta}{\sigma^t} g^t \quad \sigma^t = \sqrt{\alpha(\sigma^{t-1})^2 + (1 - \alpha)(g^t)^2}$$


Root Mean Square of the gradients  
with previous gradients being decayed

# Hard to find optimal network parameters




# In physical world .....

- Momentum





# Review: Vanilla Gradient Descent



Start at position  $\theta^0$

Compute gradient at  $\theta^0$

Move to  $\theta^1 = \theta^0 - \eta \nabla L(\theta^0)$

Compute gradient at  $\theta^1$

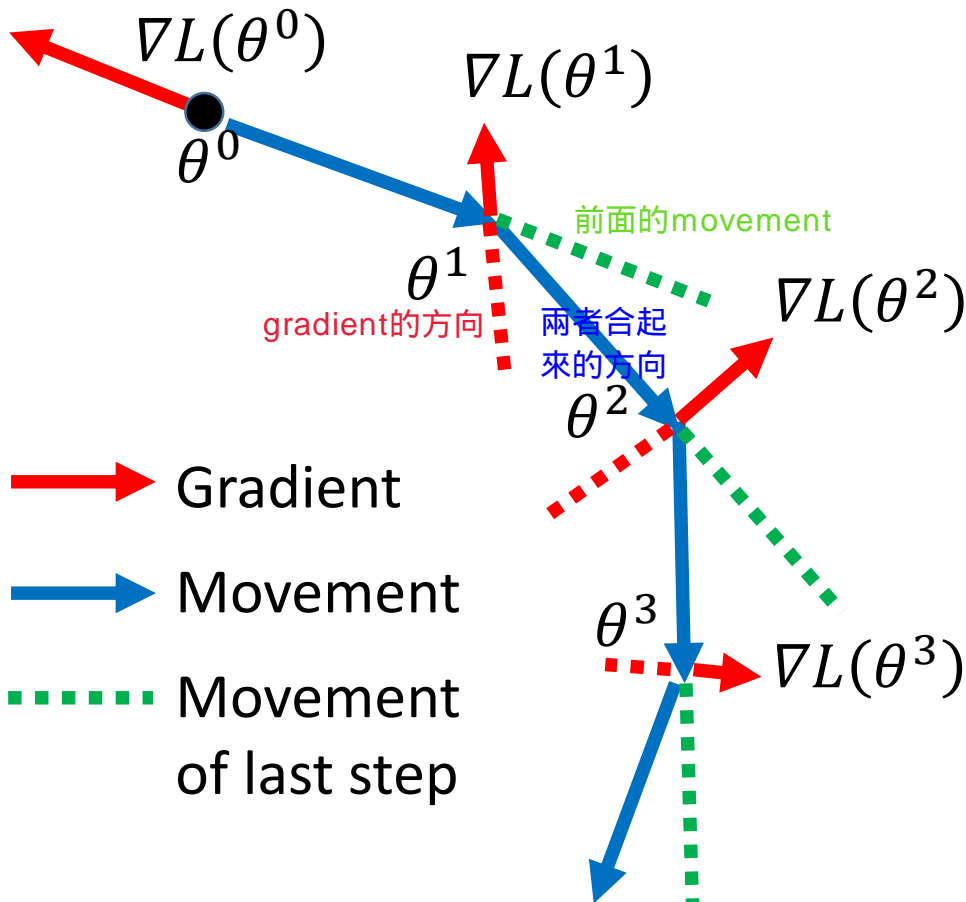
Move to  $\theta^2 = \theta^1 - \eta \nabla L(\theta^1)$

⋮

Stop until  $\nabla L(\theta^t) \approx 0$

# Momentum

Movement: movement of last step minus gradient at present



Start at point  $\theta^0$

Movement  $v^0=0$

Compute gradient at  $\theta^0$

Movement  $v^1 = \lambda v^0 - \eta \nabla L(\theta^0)$

Move to  $\theta^1 = \theta^0 + v^1$

Compute gradient at  $\theta^1$

Movement  $v^2 = \lambda v^1 - \eta \nabla L(\theta^1)$

Move to  $\theta^2 = \theta^1 + v^2$

Movement not just based on gradient, but previous movement.

# Momentum

Movement: movement of last step minus gradient at present

$v^i$  is actually the weighted sum of all the previous gradient:

$$\nabla L(\theta^0), \nabla L(\theta^1), \dots \nabla L(\theta^{i-1})$$

$$v^0 = 0$$

$$v^1 = -\eta \nabla L(\theta^0)$$

$$v^2 = -\lambda \eta \nabla L(\theta^0) - \eta \nabla L(\theta^1)$$

$\vdots$

Start at point  $\theta^0$

Movement  $v^0 = 0$

Compute gradient at  $\theta^0$

Movement  $v^1 = \lambda v^0 - \eta \nabla L(\theta^0)$

Move to  $\theta^1 = \theta^0 + v^1$

Compute gradient at  $\theta^1$


Movement  $v^2 = \lambda v^1 - \eta \nabla L(\theta^1)$

Move to  $\theta^2 = \theta^1 + v^2$

Movement not just based on gradient, but previous movement

# Momentum

Still not guarantee reaching global minima, but give some hope .....



# Adam

## RMSProp + Momentum

**Algorithm 1:** *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation.  $g_t^2$  indicates the elementwise square  $g_t \odot g_t$ . Good default settings for the tested machine learning problems are  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$  and  $\epsilon = 10^{-8}$ . All operations on vectors are element-wise. With  $\beta_1^t$  and  $\beta_2^t$  we denote  $\beta_1$  and  $\beta_2$  to the power  $t$ .

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1)$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)  $\rightarrow$  for momentum

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)  $\rightarrow$  for RMSprop

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

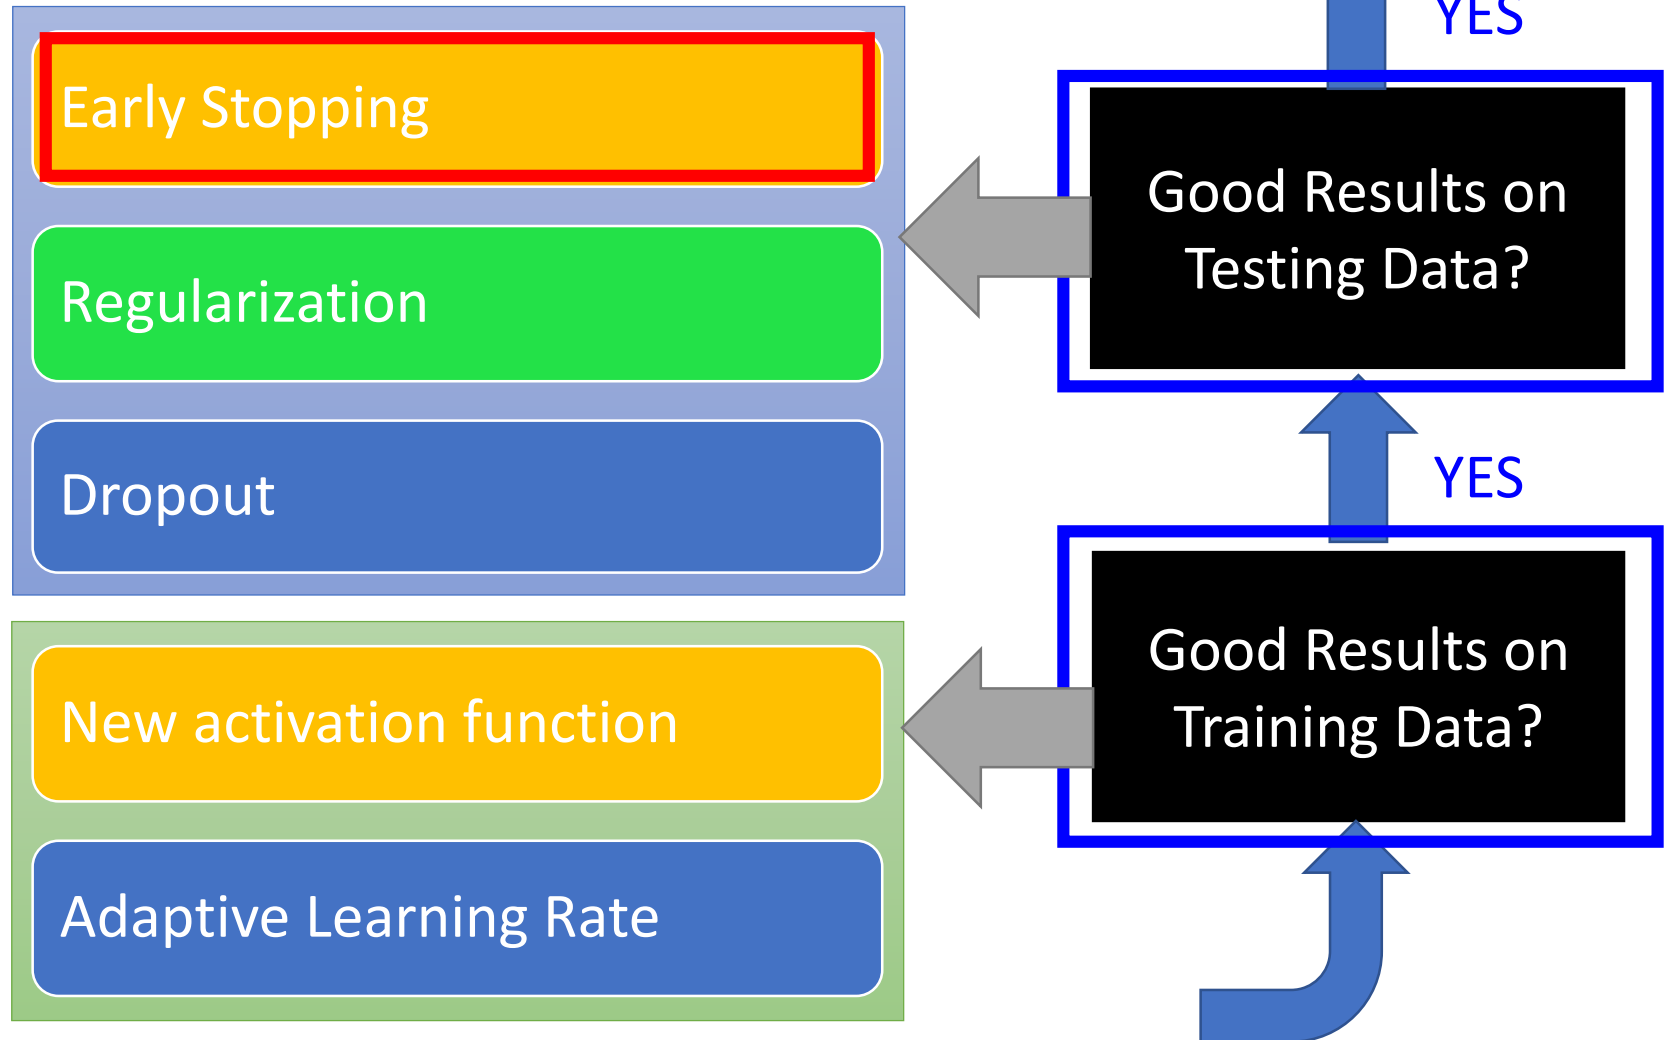
$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

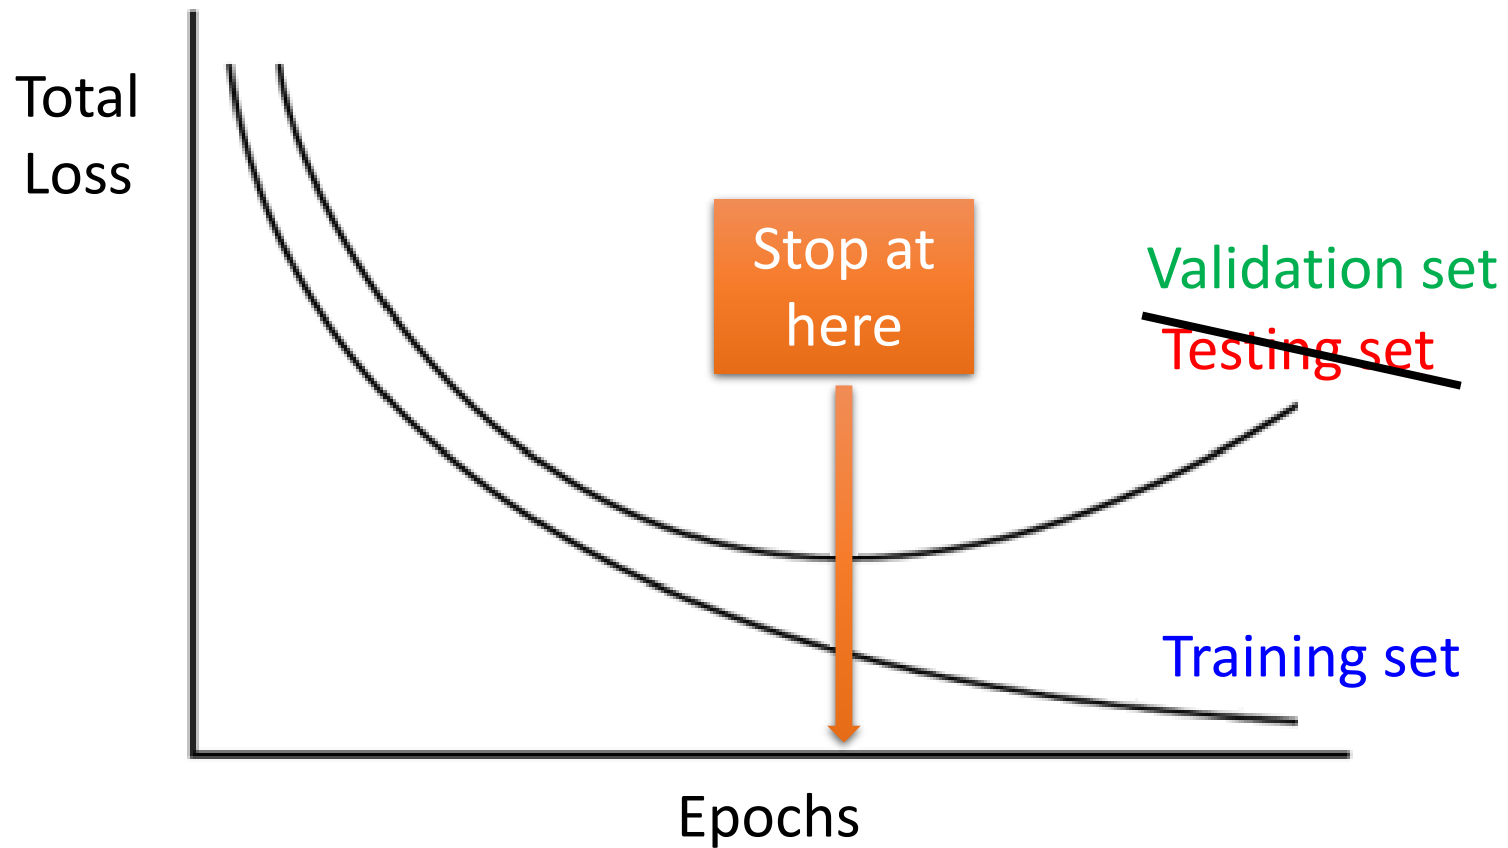
**end while**

**return**  $\theta_t$  (Resulting parameters)

# Recipe of Deep Learning

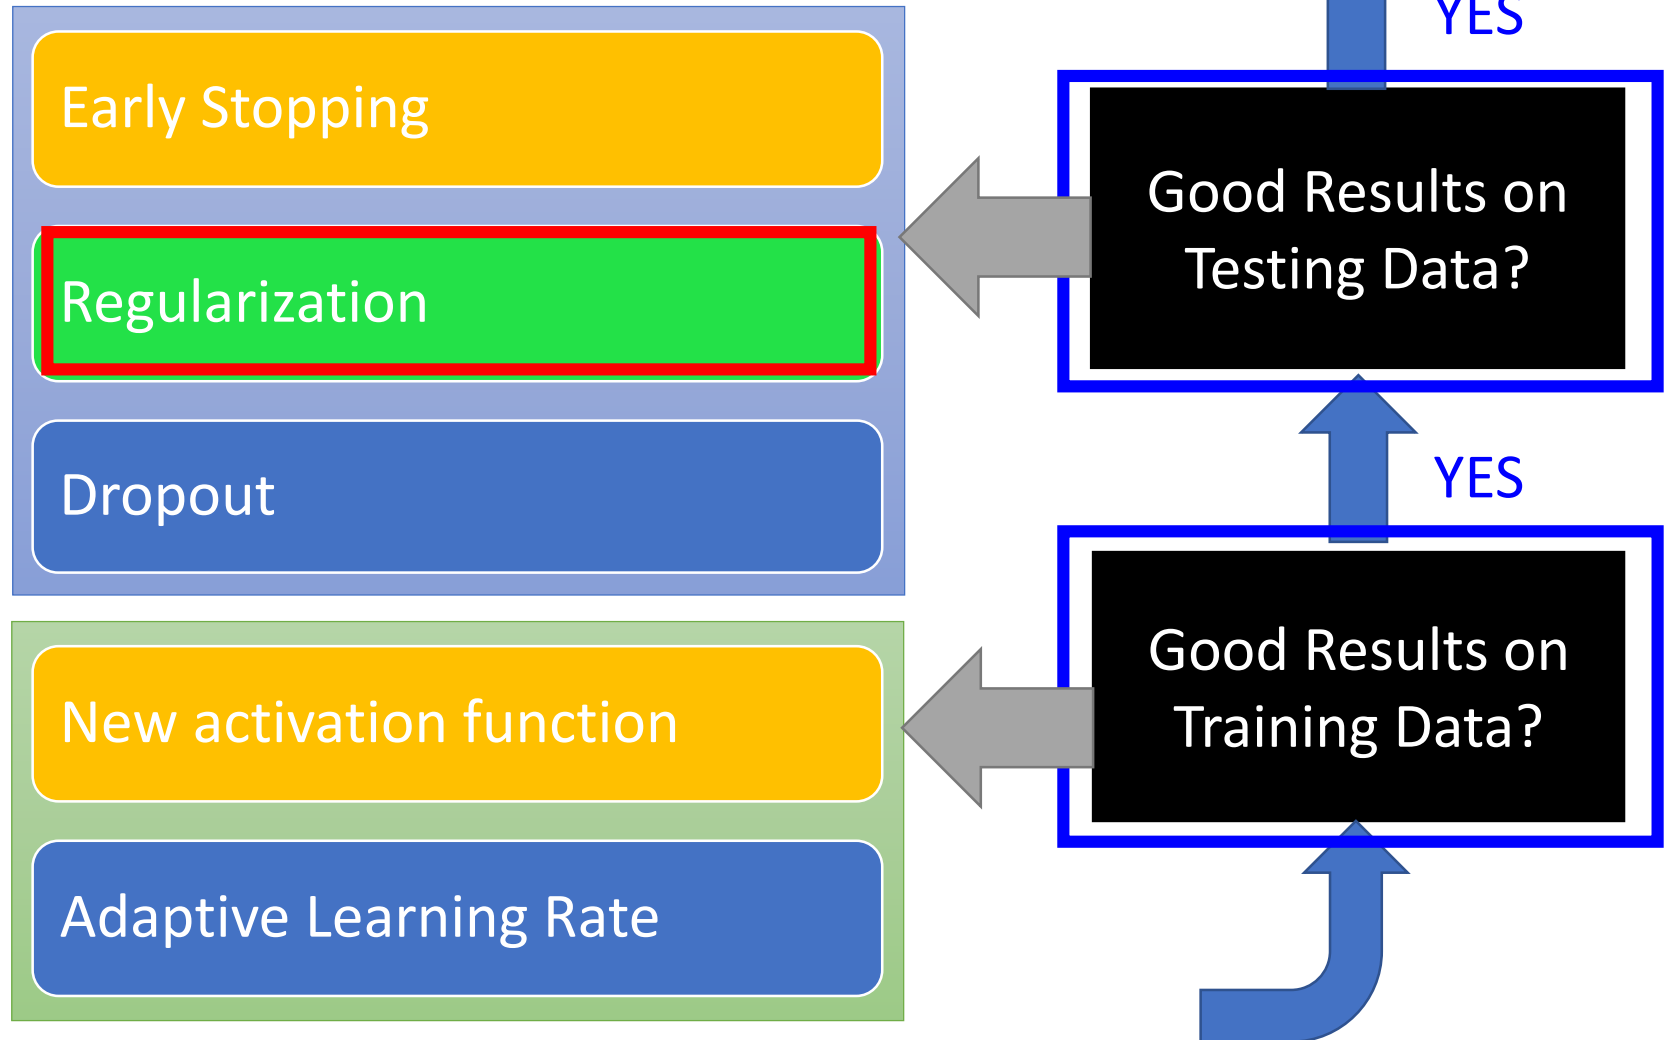


# Early Stopping



Keras: <http://keras.io/getting-started/faq/#how-can-i-interrupt-training-when-the-validation-loss-isnt-decreasing-anymore>

# Recipe of Deep Learning





# Regularization

- New loss function to be minimized
  - Find a set of weight not only minimizing original cost but also close to zero

$$L'(\theta) = \underbrace{L(\theta)} + \lambda \frac{1}{2} \underbrace{\|\theta\|_2^2} \rightarrow \text{Regularization term}$$

$$\theta = \{w_1, w_2, \dots\}$$

Original loss

(e.g. minimize square error, cross entropy ...)

L2 regularization:

$$\|\theta\|_2^2 = (w_1)^2 + (w_2)^2 + \dots$$

(usually not consider biases)

# Regularization

L2 regularization:

$$\|\theta\|_2 = (w_1)^2 + (w_2)^2 + \dots$$

- New loss function to be minimized

$$L'(\theta) = L(\theta) + \lambda \frac{1}{2} \|\theta\|_2^2 \quad \text{Gradient: } \frac{\partial L'}{\partial w} = \frac{\partial L}{\partial w} + \lambda w$$

$$\text{Update: } w^{t+1} \rightarrow w^t - \eta \frac{\partial L'}{\partial w} = w^t - \eta \left( \frac{\partial L}{\partial w} + \lambda w^t \right)$$

$$= \underbrace{(1 - \eta\lambda)w^t}_{\downarrow} - \eta \underbrace{\frac{\partial L}{\partial w}}$$

Closer to zero

Weight Decay

L1 regularization:

# Regularization

$$\|\theta\|_1 = |w_1| + |w_2| + \dots$$

- New loss function to be minimized

$$L'(\theta) = L(\theta) + \lambda \frac{1}{2} \|\theta\|_1 \quad \frac{\partial L'}{\partial w} = \frac{\partial L}{\partial w} + \lambda \operatorname{sgn}(w)$$


Update:

$$\begin{aligned} w^{t+1} &\rightarrow w^t - \eta \frac{\partial L'}{\partial w} = w^t - \eta \left( \frac{\partial L}{\partial w} + \lambda \operatorname{sgn}(w^t) \right) \\ &= w^t - \eta \frac{\partial L}{\partial w} - \underline{\eta \lambda \operatorname{sgn}(w^t)} \quad \text{Always delete} \\ &= (1 - \eta \lambda) w^t - \eta \frac{\partial L}{\partial w} \quad \text{..... L2} \end{aligned}$$

# Regularization - Weight Decay

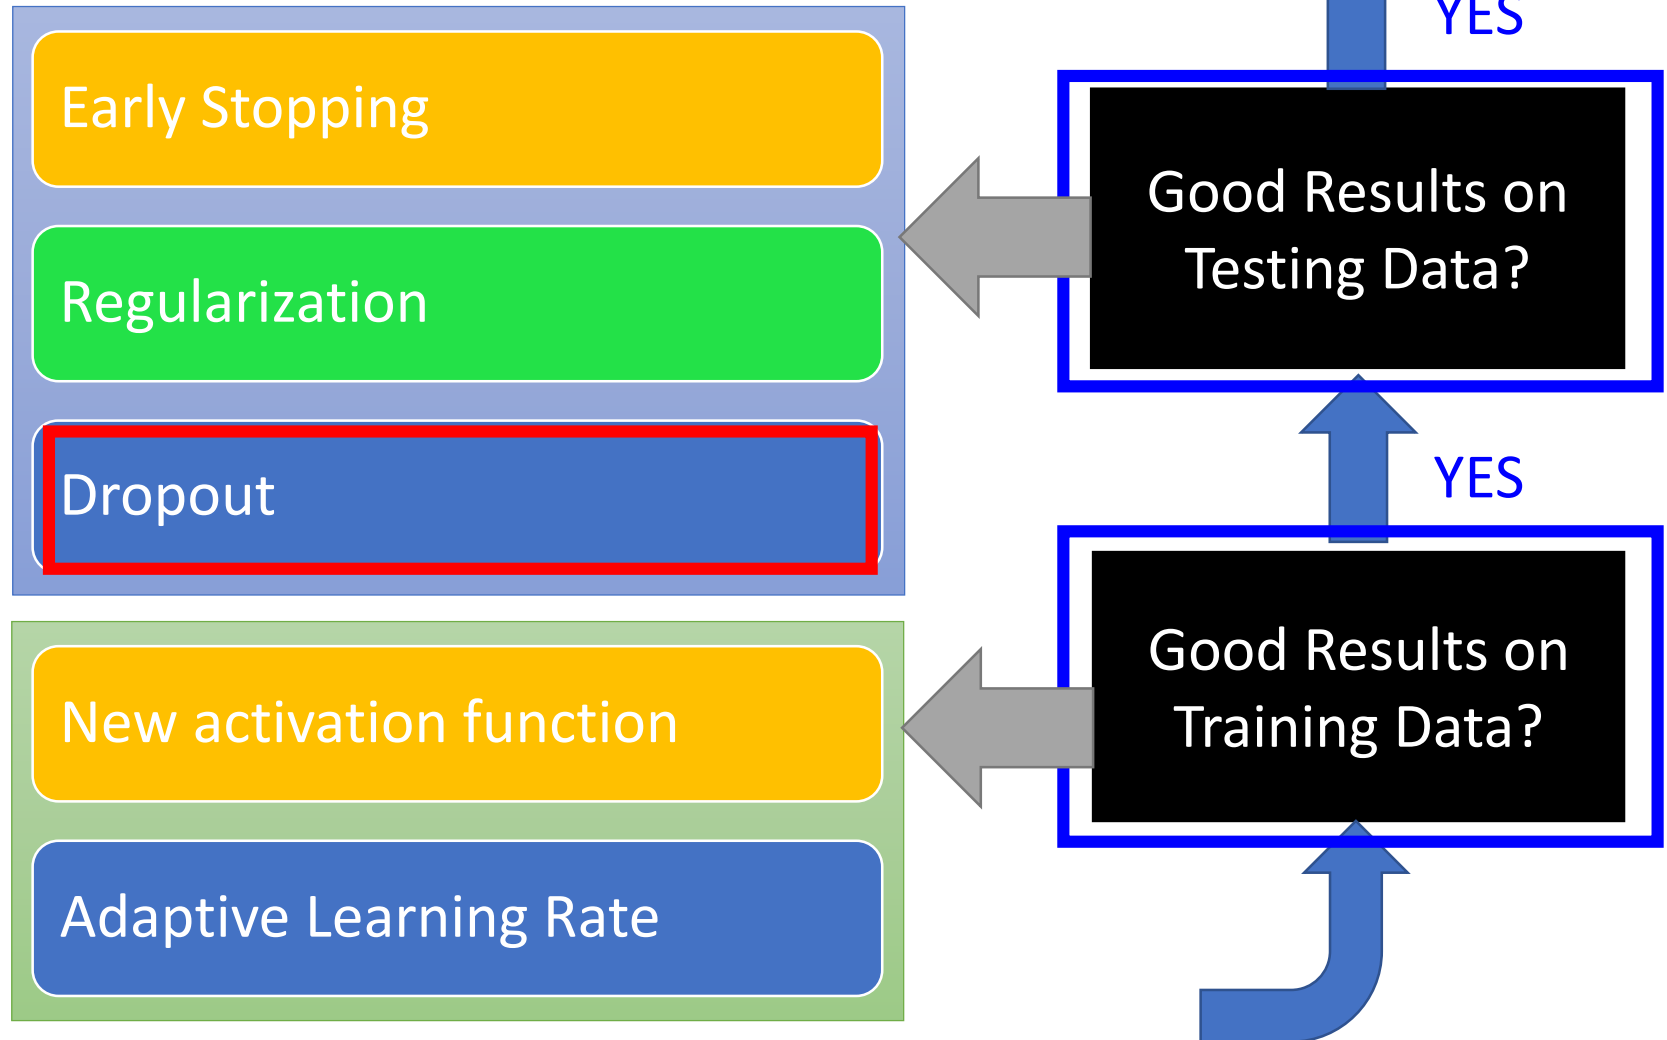
- Our brain prunes out the useless link between neurons.

沒有更新(使用)的參數會刪除




Doing the same thing to machine's brain improves the performance.

# Recipe of Deep Learning



# Dropout

Training:




- Each time before updating the parameters
  - Each neuron has  $p\%$  to dropout

丢掉X的neuron

# Dropout

Training:



➤ Each time before updating the parameters

- Each neuron has  $p\%$  to dropout




**The structure of the network is changed.**

- Using the new network for training
- 每次皆須重新sampling，故每次的 structure 都不同

For each mini-batch, we resample the dropout neurons

# Dropout

Testing:



## ➤ No dropout

- If the dropout rate at training is  $p\%$ , all the weights times  $1-p\%$
- Assume that the dropout rate is 50%.  
If a weight  $w = 1$  by training, set  $w = 0.5$  for testing.



# Dropout

## - Intuitive Reason

### Training

Dropout (腳上綁重物)




### Testing

No dropout  
(拿下重物後就變很強)



# Dropout - Intuitive Reason




- When teams up, if everyone expect the partner will do the work, nothing will be done finally.
- However, if you know your partner will dropout, you will do better.
- When testing, no one dropout actually, so obtaining good results eventually.

# Dropout - Intuitive Reason

- Why the weights should multiply  $(1-p)\%$  (dropout rate) when testing?


## Training of Dropout

Assume dropout rate is 50%




## Testing of Dropout

No dropout




# Dropout is a kind of ensemble.




Train a bunch of networks with different structures

# Dropout is a kind of ensemble.

## Ensemble




# Dropout is a kind of ensemble.









- Using one mini-batch to train one network
- Some parameters in the network are shared

# Dropout is a kind of ensemble.

## Testing of Dropout

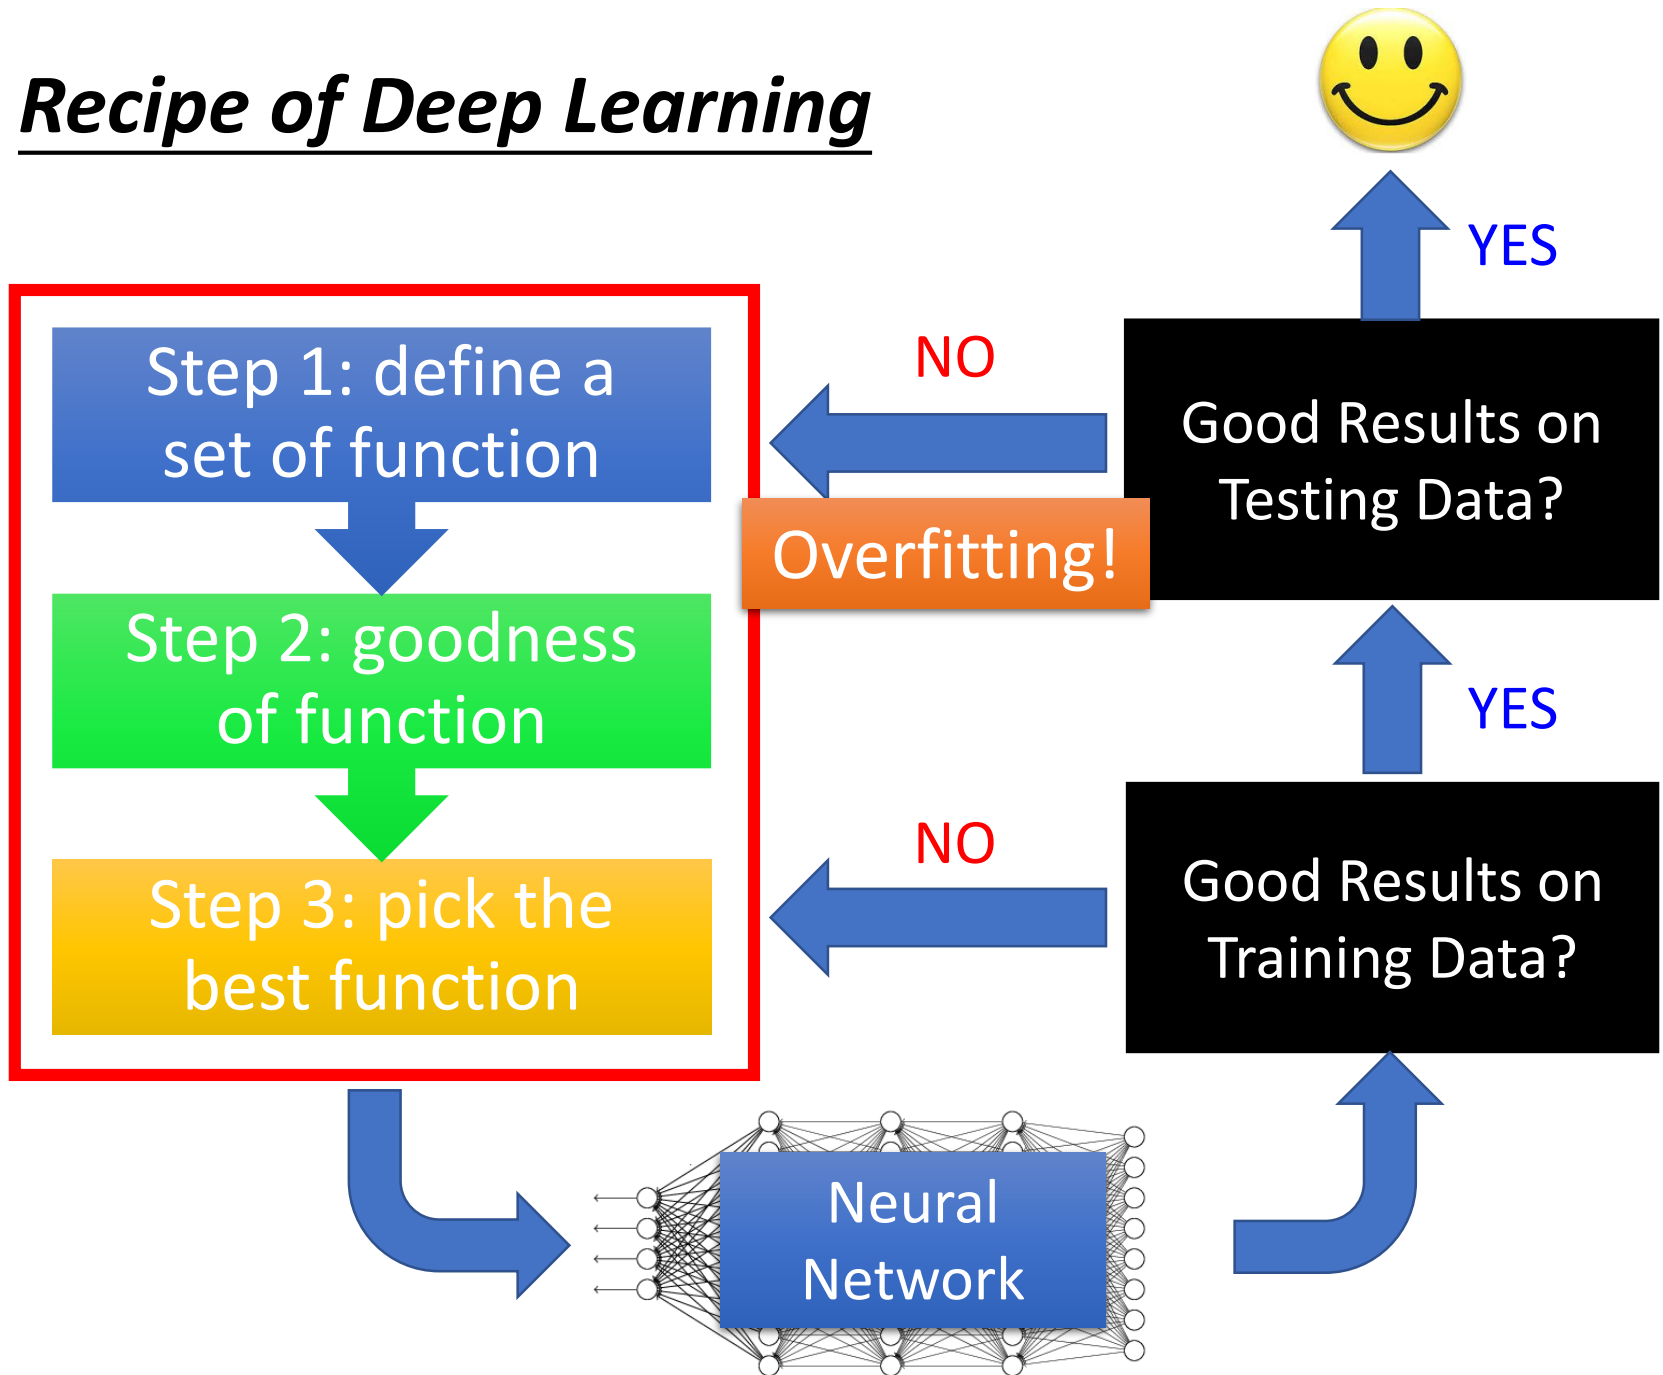


# Testing of Dropout





# Recipe of Deep Learning



# Try another task

“stock” in document

Machine

政治

經濟

體育

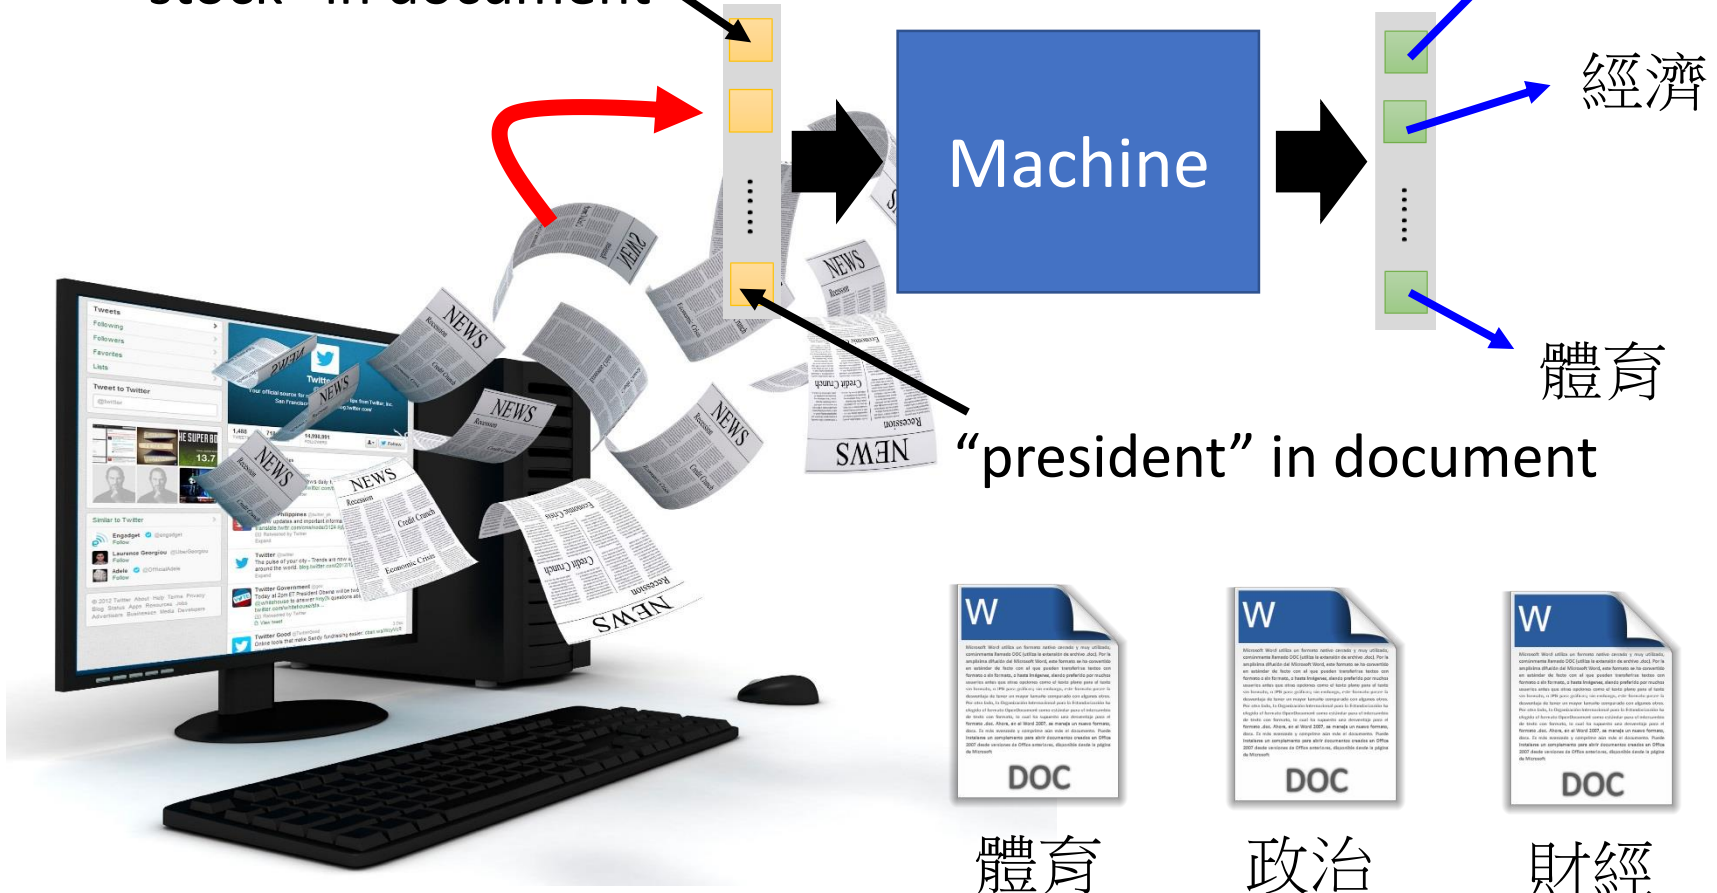
“president” in document

體育

政治

財經

<http://top-breaking-news.com/>



# Try another task

```
In [8]: x_train.shape
Out[8]: (8982, 1000)
```

```
In [9]: y_train.shape
Out[9]: (8982, 46)
```

```
In [12]: x_train[0]
```

Out [12] :

```
array([ 0.,  1.,  1.,  0.,  1.,  1.,  1.,  1.,  1. Out[10]: (2246, 1000)
        0.,  0.,  1.,  1.,  1.,  0.,  1.,  0.,  0
        1.,  0.,  0.,  1.,  1.,  0.,  1.,  0.,  0
        1.,  0.,  0.,  0.,  1.,  1.,  0.,  0.,  0
        1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0
        0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  0.,  0.,  0.,  1.,  1.,  0.,  0.,  0.,  0.,  1.,  1.,
        0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,
        0.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  1.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,  1.,  1.,
        0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,
        0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
        0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,
```

```
In [13]: y_train[0]
```

Out[13]:

[illegible]

# Live Demo